

ANSIBLE DYNAMIC ASSIGNMENTS (INCLUDE) AND COMMUNITY ROLES

Last 2 projects have already equipped us with some knowledge and skills on Ansible, so we can perform configurations using *playbooks*, *roles* and *imports*. Now we will continue configuring our UAT servers, learning and practicing new Ansible concepts and modules.

In this project we will introduce *dynamic assignments* by using *include* module.

Now you may be wondering, what is the difference between **static** and **dynamic** assignments?

Well, from *Project 12*, you can already tell that static assignments use *import* Ansible module. The module that enables dynamic assignments is *include*.

Hence,

```
import = Static
include = Dynamic
```

When the *import* module is used, all statements are pre-processed at the time playbooks are *parsed*. Meaning, when you execute *site.yml* playbook, Ansible will process all the playbooks referenced during the time it is parsing the statements. This also means that, during actual execution, if any statement changes, such statements will not be considered. Hence, it is static.

On the other hand, when *include* module is used, all statements are processed only during execution of the playbook. Meaning, after the statements are *parsed*, any changes to the statements encountered during execution will be used.

INTRODUCING DYNAMIC ASSIGNMENT INTO OUR STRUCTURE

In your <https://github.com/<your-name>/ansible-config-mgt> GitHub repository start a new branch and call it **dynamic-assignments**.

Create a new folder, name it **dynamic-assignments**. Then inside this folder, create a new file and name it *env-vars.yml*

```

ubuntu@ip-172-31-37-181:~/ansible-config-mgt$ git checkout dynamic-assignments
Branch 'dynamic-assignments' set up to track remote branch 'dynamic-assignments' from 'origin'.
Switched to a new branch 'dynamic-assignments'
ubuntu@ip-172-31-37-181:~/ansible-config-mgt$ git branch
* dynamic-assignments
  feature/proj-45
  main
  refactor
ubuntu@ip-172-31-37-181:~/ansible-config-mgt$ mkdir dynamic-assignments
ubuntu@ip-172-31-37-181:~/ansible-config-mgt$ ls
Ansible.md  dynamic-assignments  inventory  playbooks  roles  static-assignments
ubuntu@ip-172-31-37-181:~/ansible-config-mgt$ cd dynamic-assignments/
ubuntu@ip-172-31-37-181:~/ansible-config-mgt/dynamic-assignments$ sudo vi env-vars.yml
ubuntu@ip-172-31-37-181:~/ansible-config-mgt/dynamic-assignments$ ls
ubuntu@ip-172-31-37-181:~/ansible-config-mgt/dynamic-assignments$ sudo touch env-vars.yml
ubuntu@ip-172-31-37-181:~/ansible-config-mgt/dynamic-assignments$ ls
env-vars.yml
ubuntu@ip-172-31-37-181:~/ansible-config-mgt/dynamic-assignments$

```

Note: Depending on what method you used in the previous project you may have or not have `roles` folder in your GitHub repository - if you used `ansible-galaxy`, then `roles` directory was only created on your `Jenkins-Ansible` server locally. It is recommended to have all the codes managed and tracked in GitHub, so you might want to recreate this structure manually in this case - it is up to you.

Our GitHub should have following structure by now.

```

├── dynamic-assignments
│   └── env-vars.yml
├── inventory
│   ├── dev
│   ├── stage
│   ├── uat
│   └── prod
├── playbooks
│   └── site.yml
├── roles (optional folder)
│   ├── defaults
│   ├── handlers
│   ├── meta
│   ├── tasks
│   ├── templates
│   └── README.md
├── static-assignments
│   ├── common-del.yml
│   ├── common.yml
│   └── uat-webservers.yml

```

Since we will be using the same Ansible to configure multiple environments, and each of these environments will have certain unique attributes, such as

servername, ip-address etc., we will need a way to set values to variables per specific environment.

For this reason, we will now create a folder to keep each environment's variables file. Therefore, create a new folder **env-vars**, then for each environment, create new **YAML** files which we will use to set variables.

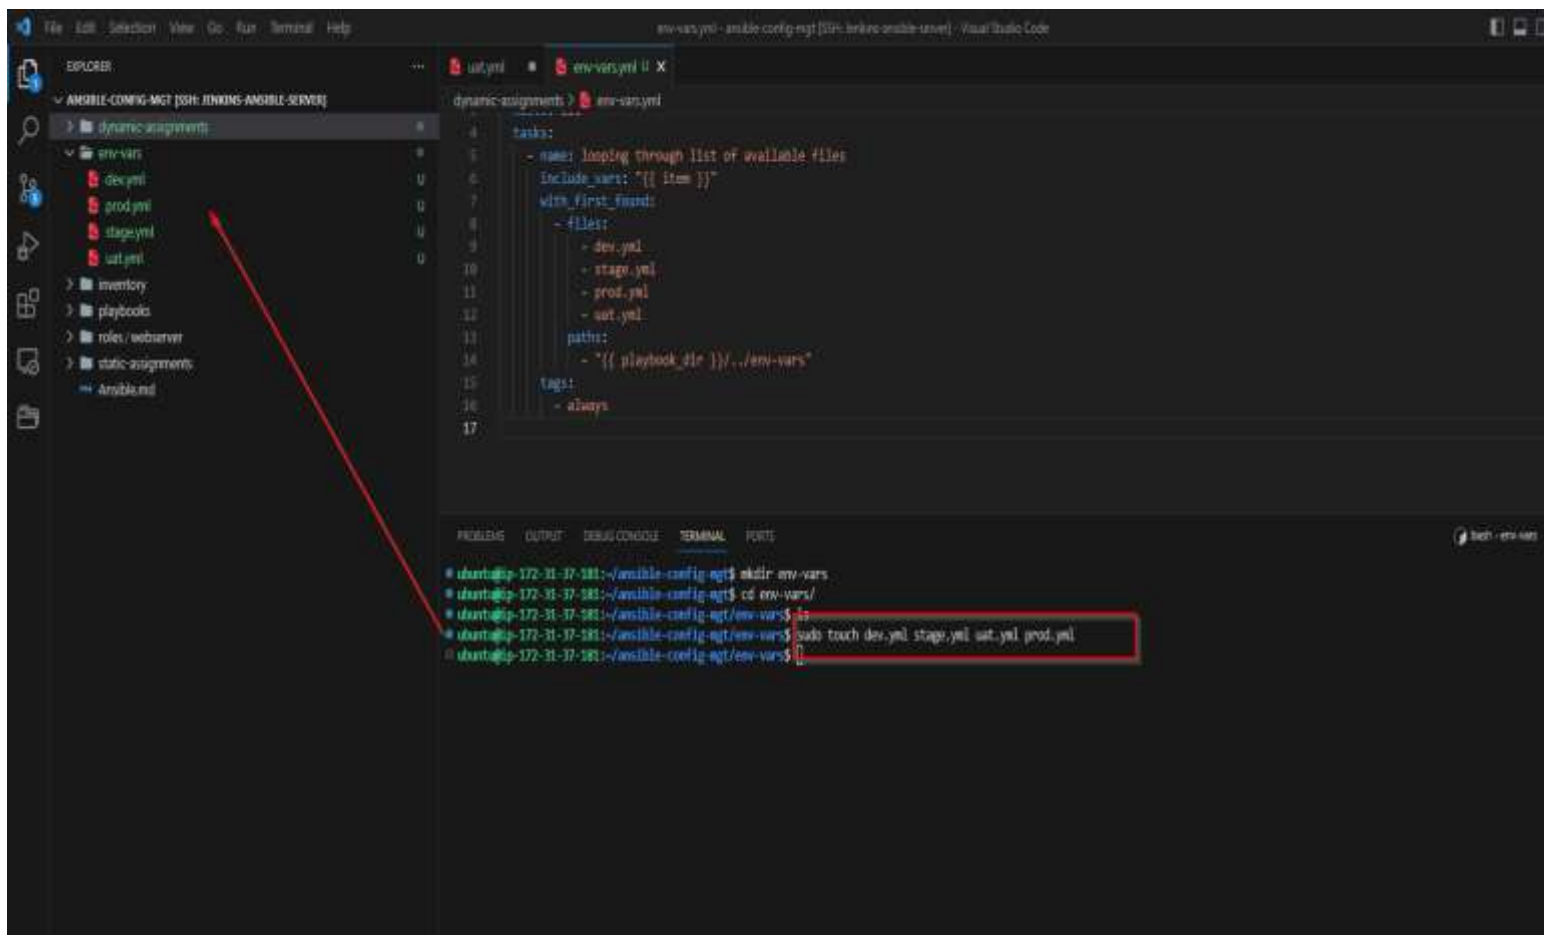
```
$ cd ansible-config-mgt
```

```
$ sudo mkdir env-vars
```

```
$ sudo touch dev.yml stage.yml uat.yml prod.yml
```

Our layout should now look like this.

```
├── dynamic-assignments
│   └── env-vars.yml
├── env-vars
│   ├── dev.yml
│   ├── stage.yml
│   ├── uat.yml
│   └── prod.yml
├── inventory
│   ├── dev
│   ├── stage
│   ├── uat
│   └── prod
├── playbooks
│   └── site.yml
└── static-assignments
    ├── common.yml
    └── webservers.yml
```



Now paste the instruction below into the *env-vars.yml* file.

```
---
- name: collate variables from env specific file, if it exists
  hosts: all
  tasks:
    - name: looping through list of available files
      include_vars: "{{ item }}"
      with_first_found:
        - files:
            - dev.yml
            - stage.yml
            - prod.yml
            - uat.yml
        paths:
            - "{{ playbook_dir }}/../env-vars"
      tags:
        - always
```

```
ubuntu@ip-172-31-37-181:~/ansible-config-mgt/dynamic-assignments$ sudo vi env-vars.yml
ubuntu@ip-172-31-37-181:~/ansible-config-mgt/dynamic-assignments$ ls
env-vars.yml
ubuntu@ip-172-31-37-181:~/ansible-config-mgt/dynamic-assignments$ cat env-vars.yml
```

```
- name: collate variables from env specific file, if it exists
  hosts: all
  tasks:
    - name: looping through list of available files
      include_vars: "{{ item }}"
      with_first_found:
        - files:
            - dev.yml
            - stage.yml
            - prod.yml
            - uat.yml
          paths:
            - "{{ playbook_dir }}/../env-vars"
      tags:
        - always
```

```
ubuntu@ip-172-31-37-181:~/ansible-config-mgt/dynamic-assignments$
```

Notice 3 things to notice here:

1. We used `include_vars` syntax instead of `include`, this is because Ansible developers decided to separate different features of the module. From Ansible version **2.8**, the `include` module is deprecated and variants of `include_*` must be used. These are:
 - [include_role](#)
 - [include_tasks](#)
 - [include_vars](#)

In the same version, variants of **import** were also introduced, such as:

- [import_role](#)
 - [import_tasks](#)
2. We made use of a [special variables](#) `{ playbook_dir }` and `{ inventory_file }`. `{ playbook_dir }` will help Ansible to determine the location of the running playbook, and from there navigate to other path on the filesystem. `{ inventory_file }` on the other hand will dynamically resolve to the name of the inventory file being used, then append `.yaml` so that it picks up the required file within the `env-vars` folder.
 3. We are including the variables using a loop. `with_first_found` implies that, looping through the list of files, the first one found is used. This is good so that we can always set default values in case an environment specific env file does not exist.

UPDATE SITE.YML WITH DYNAMIC ASSIGNMENTS

We will now update `site.yml` file to make use of the dynamic assignment. (At this point, we cannot test it yet. We are just setting the stage for what is yet to come)

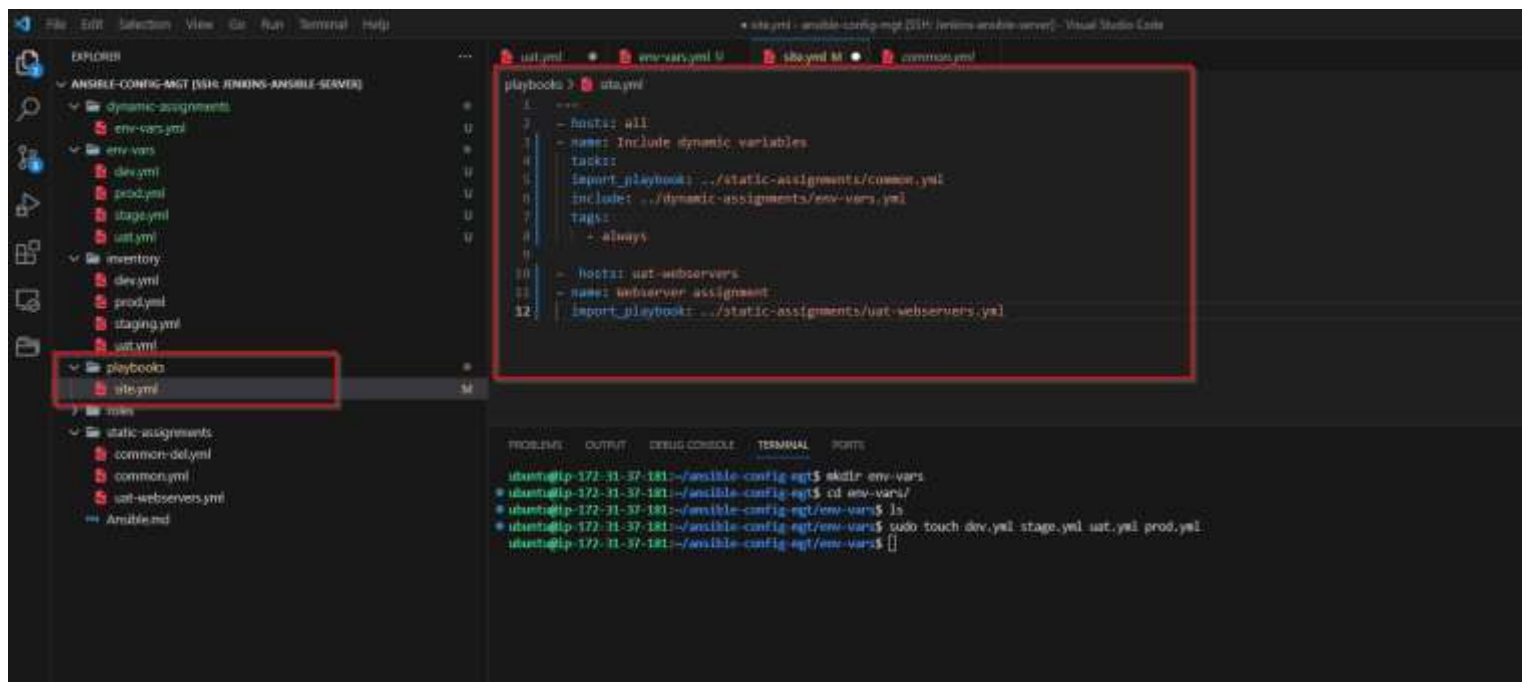
Update the `site.yml` to look like below:

```

---
- hosts: all
- name: Include dynamic variables
  tasks:
  import_playbook: ../static-assignments/common.yml
  include: ../dynamic-assignments/env-vars.yml
  tags:
    - always

- hosts: uat-webservers
- name: Webserver assignment
  import_playbook: ../static-assignments/uat-ebservers.yml

```



COMMUNITY ROLES

Before we proceed, let us commit and push new changes on **dynamic-assignment** branch to git.


```

ubuntu@ip-172-31-37-181:~/ansible-config-mgt$ git status
On branch dynamic-assignments
Your branch is up to date with 'origin/dynamic-assignments'.

Changes not staged for commit:
  (use "git add <file>..." to update what will be committed)
  (use "git restore <file>..." to discard changes in working directory)
        modified:   playbooks/site.yml

Untracked files:
  (use "git add <file>..." to include in what will be committed)
        dynamic-assignments/
        env-vars/

no changes added to commit (use "git add" and/or "git commit -a")
ubuntu@ip-172-31-37-181:~/ansible-config-mgt$ git add .
ubuntu@ip-172-31-37-181:~/ansible-config-mgt$ git commit -m "adding vars"
[dynamic-assignments ad78cf8] adding vars
6 files changed, 25 insertions(+), 3 deletions(-)
create mode 100644 dynamic-assignments/env-vars.yml
create mode 100644 env-vars/dev.yml
create mode 100644 env-vars/prod.yml
create mode 100644 env-vars/stage.yml
create mode 100644 env-vars/uat.yml

```

```

ubuntu@ip-172-31-37-181:~/ansible-config-mgt/env-vars$ git push
Enumerating objects: 10, done.
Counting objects: 100% (10/10), done.
Compressing objects: 100% (5/5), done.
Writing objects: 100% (7/7), 988 bytes | 329.00 KiB/s, done.
Total 7 (delta 0), reused 0 (delta 0)
To https://github.com/obaigbenaa/ansible-config-mgt.git
df83a6e..ad78cf8 dynamic-assignments -> dynamic-assignments

```

Now it is time to create a role for MySQL database – it should install the MySQL package, create a database and configure users. But why should we re-invent the wheel? There are tons of roles that have already been developed by other open source engineers out there. These roles are actually production ready, and dynamic to accommodate most of Linux flavours. With Ansible Galaxy again, we can simply download a ready to use ansible role, and move on with other businesses.

Download Mysql Ansible Role

We can browse available community roles [here](#)

We will be using a *MySQL role developed by geerlingguy*.

Hint: To preserve your your GitHub in actual state after you install a new role – make a commit and push to master your ‘ansible-config-mgt’ directory. Of course you must have **git** installed and configured on **Jenkins-Ansible** server and, for more

convenient work with codes, you can configure Visual Studio Code to work with this directory.

We will now create a new branch called '**roles-feature**' and move into it. This is where we will be creating the role for MySQL database.

```
$ git branch roles-feature
```

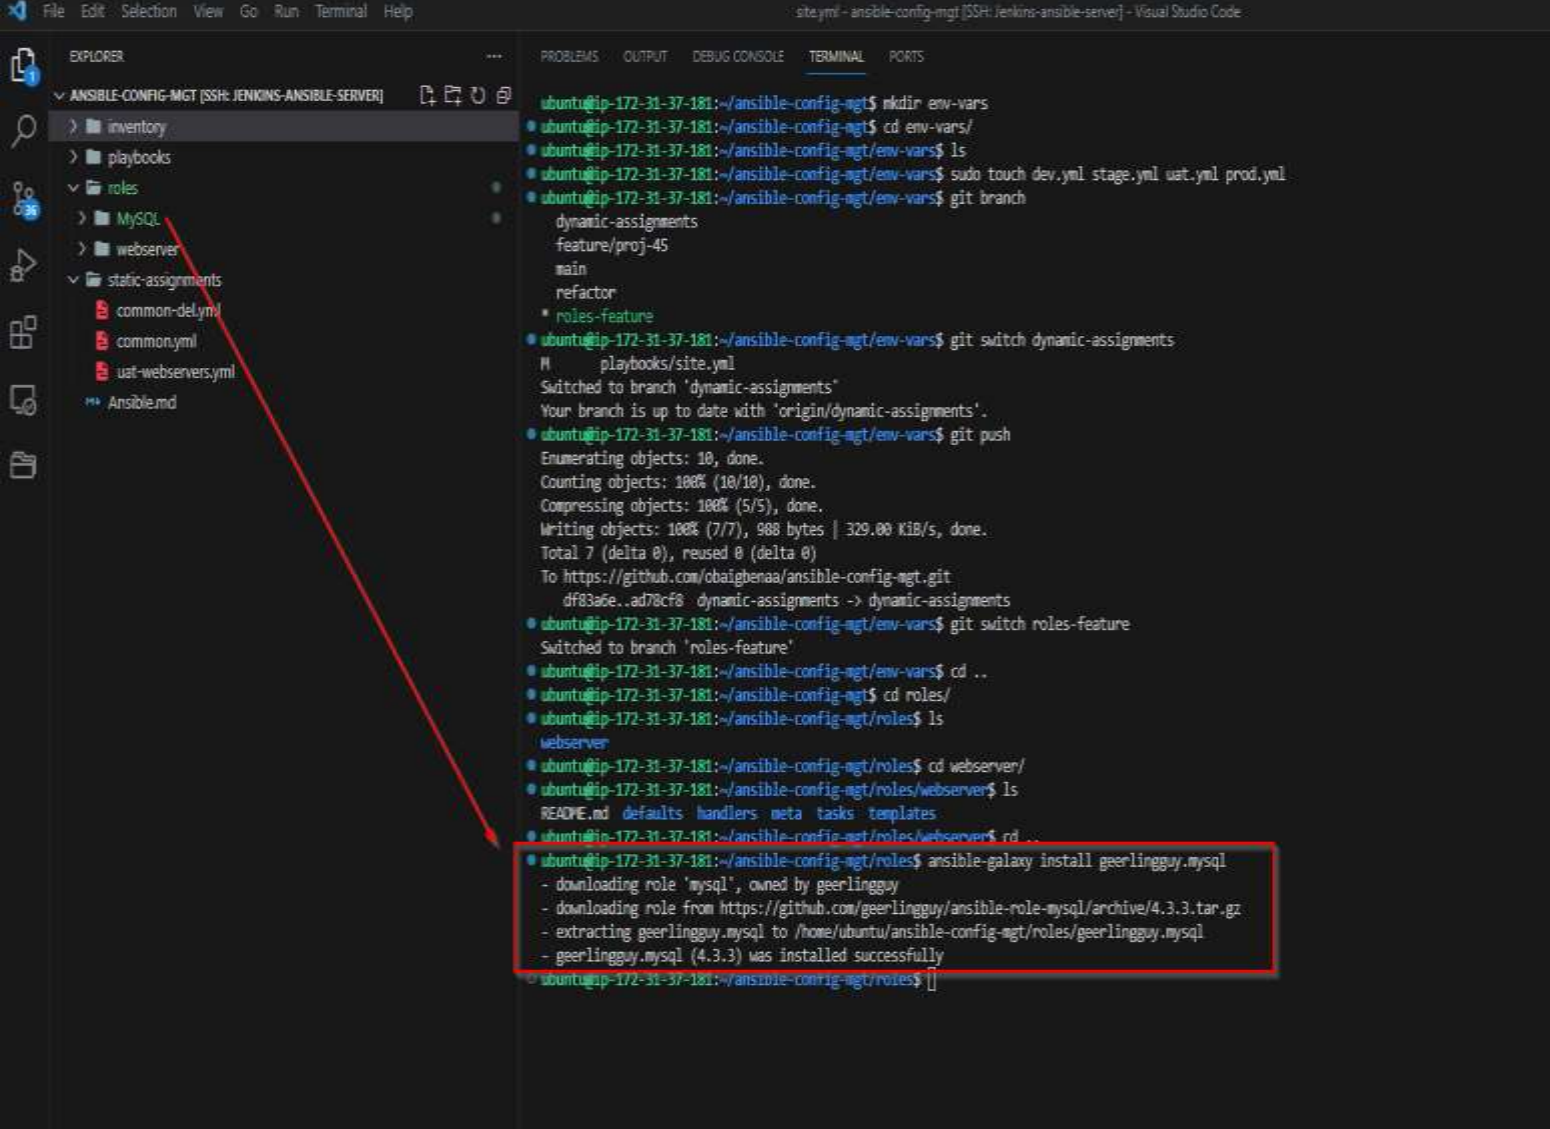
```
$ git switch roles-feature
```

Now we are in the **roles-feature** branch.

Inside **roles** directory we will now create a new **MySQL** role with **ansible-galaxy** install **geerlingguy.mysql** and rename the folder to **mysql**

```
$ ansible-galaxy install geerlingguy.mysql
```

```
$ mv geerlingguy.mysql/ mysql
```



```
File Edit Selection View Go Run Terminal Help
site.yml - ansible-config-mgt [SSH: Jenkins-Ansible-Server] - Visual Studio Code

EXPLORER
ANSIBLE-CONFIG-MGT [SSH: JENKINS-ANSIBLE-SERVER]
  inventory
  playbooks
  roles
  MySQL
  webserver
  static-assignments
  common-del.yml
  common.yml
  uat-webservers.yml
  Ansible.md

PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL PORTS

ubuntu@ip-172-31-37-181:~/ansible-config-mgt$ mkdir env-vars
ubuntu@ip-172-31-37-181:~/ansible-config-mgt$ cd env-vars/
ubuntu@ip-172-31-37-181:~/ansible-config-mgt/env-vars$ ls
ubuntu@ip-172-31-37-181:~/ansible-config-mgt/env-vars$ sudo touch dev.yml stage.yml uat.yml prod.yml
ubuntu@ip-172-31-37-181:~/ansible-config-mgt/env-vars$ git branch
dynamic-assignments
feature/proj-45
main
refactor
* roles-feature
ubuntu@ip-172-31-37-181:~/ansible-config-mgt/env-vars$ git switch dynamic-assignments
M      playbooks/site.yml
Switched to branch 'dynamic-assignments'
Your branch is up to date with 'origin/dynamic-assignments'.
ubuntu@ip-172-31-37-181:~/ansible-config-mgt/env-vars$ git push
Enumerating objects: 10, done.
Counting objects: 100% (10/10), done.
Compressing objects: 100% (5/5), done.
Writing objects: 100% (7/7), 988 bytes | 329.00 KiB/s, done.
Total 7 (delta 0), reused 0 (delta 0)
To https://github.com:obaigbenaa/ansible-config-mgt.git
df83a6e..ad78cf8  dynamic-assignments -> dynamic-assignments
ubuntu@ip-172-31-37-181:~/ansible-config-mgt/env-vars$ git switch roles-feature
Switched to branch 'roles-feature'
ubuntu@ip-172-31-37-181:~/ansible-config-mgt/env-vars$ cd ..
ubuntu@ip-172-31-37-181:~/ansible-config-mgt$ cd roles/
ubuntu@ip-172-31-37-181:~/ansible-config-mgt/roles$ ls
webserver
ubuntu@ip-172-31-37-181:~/ansible-config-mgt/roles$ cd webserver/
ubuntu@ip-172-31-37-181:~/ansible-config-mgt/roles/webserver$ ls
README.md defaults handlers meta tasks templates
ubuntu@ip-172-31-37-181:~/ansible-config-mgt/roles/webserver$ cd ..
ubuntu@ip-172-31-37-181:~/ansible-config-mgt/roles$ ansible-galaxy install geerlingguy.mysql
- downloading role 'mysql', owned by geerlingguy
- downloading role from https://github.com/geerlingguy/ansible-role-mysql/archive/4.3.3.tar.gz
- extracting geerlingguy.mysql to /home/ubuntu/ansible-config-mgt/roles/geerlingguy.mysql
- geerlingguy.mysql (4.3.3) was installed successfully
ubuntu@ip-172-31-37-181:~/ansible-config-mgt/roles$
```

Read and check the **README.md** file for instructions and edit roles configuration to use correct credentials for MySQL required for the **tooling** website.

If you scroll to 'mysql_users: []' tab of the .md file, you will find the default setup of the configuration as below:

Update it with the correct credentials and privileges to want for username, host list, password, and database access.

```
97
98 # Logging settings.
99 mysql_log: ""
100 # The following variables have a default value depending on operating system.
101 # mysql_log_error: /var/log/mysql/mysql.err
102 # mysql_syslog_tag: mysql
103
104 mysql_config_include_files: []
105 # - src: path/relative/to/playbook/file.cnf
106 # - { src: path/relative/to/playbook/anotherfile.cnf, force: yes }
107
108 # Databases.
109 mysql_databases: []
110 | - name: tooling
111 |   collation: utf8_general_ci
112 |   encoding: utf8
113 |   replicate: 1
114
115 # Users.
116 mysql_users: []
117 | - name: webaccess
118 |   host: 0.0.0.0
119 |   password: password
120 |   priv: ' *.*:ALL,GRANT'
121
122 mysql_disable_log_bin: false
123
124 # Replication settings (replication is only enabled if master/user have values).
125 mysql_server_id: "1"
126 mysql_max_binlog_size: "100M"
127 mysql_binlog_format: "ROW"
128 mysql_expire_logs_days: "10"
129 mysql_replication_role: ''
130 mysql_replication_master: ''
131 mysql_replication_master_inventory_host: "{{ mysql_replication_master }}"
```

Save the updated configuration.

Now we will upload all code changes (on branches **roles-feature** and **dynamic-assignment**) into your GitHub:

```
$ git add .
```

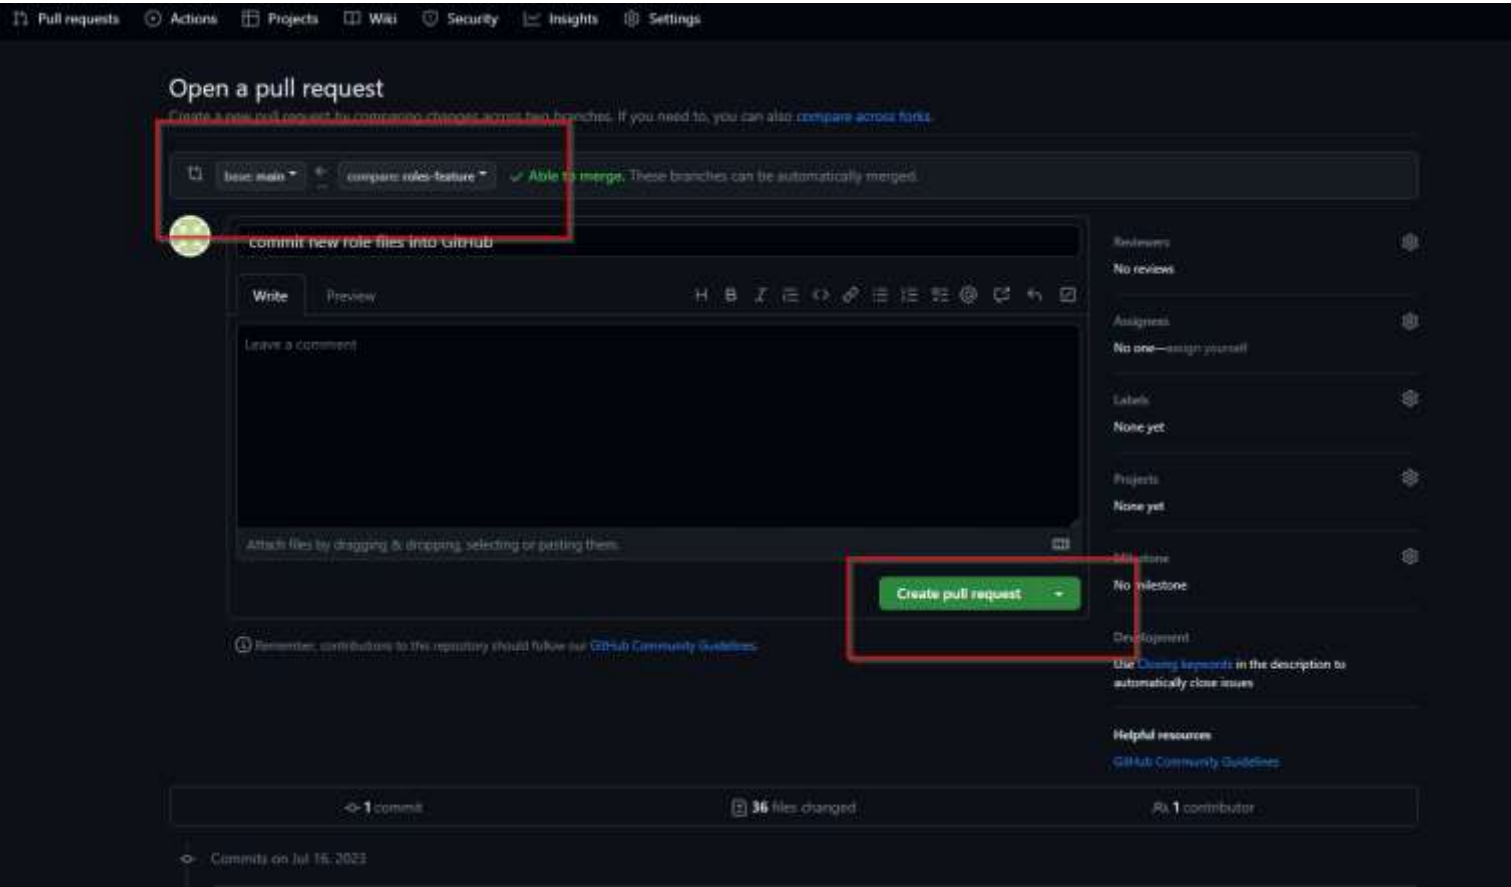
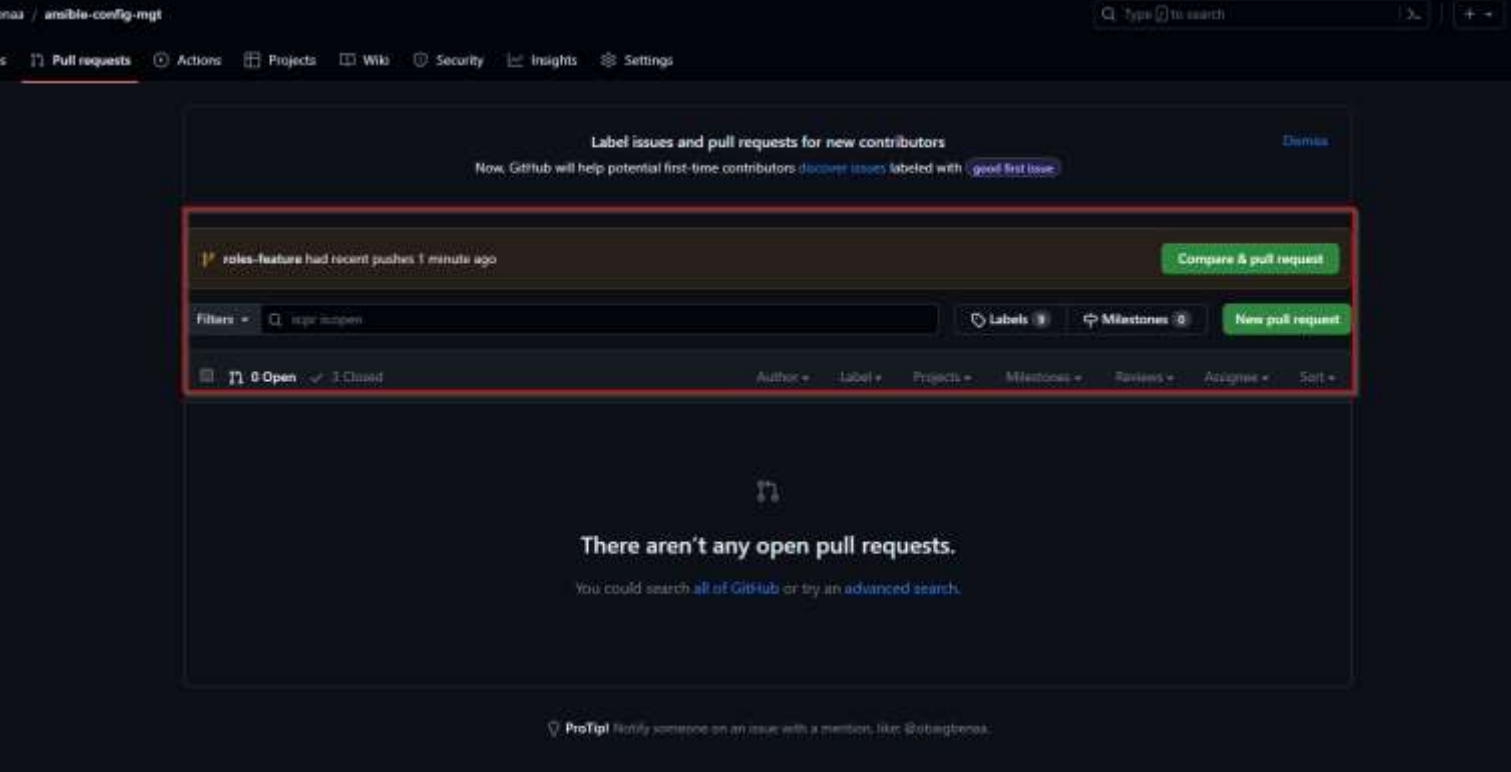
```
$ git commit -m "Commit new role files into GitHub"
```

```
$ git push --set-upstream origin roles-feature
```

PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL PORTS

```
● ubuntu@ip-172-31-37-181:~/ansible-config-mgt$ git commit -m "commit new role files into GitHub"
[roles-feature 9bf4779] commit new role files into GitHub
36 files changed, 1328 insertions(+)
create mode 100644 roles/MySQL/.ansible-lint
create mode 100644 roles/MySQL/.github/FUNDING.yml
create mode 100644 roles/MySQL/.github/stale.yml
create mode 100644 roles/MySQL/.github/workflows/ci.yml
create mode 100644 roles/MySQL/.github/workflows/release.yml
create mode 100644 roles/MySQL/.gitignore
create mode 100644 roles/MySQL/.yamllint
create mode 100644 roles/MySQL/LICENSE
create mode 100644 roles/MySQL/README.md
create mode 100644 roles/MySQL/defaults/main.yml
create mode 100644 roles/MySQL/handlers/main.yml
create mode 100644 roles/MySQL/meta/.galaxy_install_info
create mode 100644 roles/MySQL/meta/main.yml
create mode 100644 roles/MySQL/molecule/default/converge.yml
create mode 100644 roles/MySQL/molecule/default/molecule.yml
create mode 100644 roles/MySQL/tasks/configure.yml
create mode 100644 roles/MySQL/tasks/databases.yml
create mode 100644 roles/MySQL/tasks/main.yml
create mode 100644 roles/MySQL/tasks/replication.yml
create mode 100644 roles/MySQL/tasks/secure-installation.yml
create mode 100644 roles/MySQL/tasks/setup-Archlinux.yml
create mode 100644 roles/MySQL/tasks/setup-Debian.yml
create mode 100644 roles/MySQL/tasks/setup-RedHat.yml
create mode 100644 roles/MySQL/tasks/users.yml
create mode 100644 roles/MySQL/tasks/variables.yml
create mode 100644 roles/MySQL/templates/my.cnf.j2
create mode 100644 roles/MySQL/templates/root-my.cnf.j2
create mode 100644 roles/MySQL/templates/user-my.cnf.j2
create mode 100644 roles/MySQL/vars/Archlinux.yml
create mode 100644 roles/MySQL/vars/Debian-10.yml
create mode 100644 roles/MySQL/vars/Debian-11.yml
create mode 100644 roles/MySQL/vars/Debian-12.yml
create mode 100644 roles/MySQL/vars/Debian.yml
create mode 100644 roles/MySQL/vars/RedHat-7.yml
create mode 100644 roles/MySQL/vars/RedHat-8.yml
create mode 100644 roles/MySQL/vars/RedHat-9.yml
● ubuntu@ip-172-31-37-181:~/ansible-config-mgt$
```

```
create mode 100644 roles/MySQL/vars/RedHat-8.yml
create mode 100644 roles/MySQL/vars/RedHat-9.yml
● ubuntu@ip-172-31-37-181:~/ansible-config-mgt$ git push origin roles-feature
Enumerating objects: 51, done.
Counting objects: 100% (51/51), done.
Compressing objects: 100% (44/44), done.
Writing objects: 100% (49/49), 17.42 KiB | 495.00 KiB/s, done.
Total 49 (delta 4), reused 0 (delta 0)
remote: Resolving deltas: 100% (4/4), completed with 1 local object.
remote:
remote: Create a pull request for 'roles-feature' on GitHub by visiting:
remote:   https://github.com/obaigbenaa/ansible-config-mgt/pull/new/roles-feature
remote:
To https://github.com/obaigbenaa/ansible-config-mgt.git
* [new branch]   roles-feature -> roles-feature
● ubuntu@ip-172-31-37-181:~/ansible-config-mgt$
```



commit new role files into GitHub #4

Edit <> Code

Open obaigbenaa wants to merge 1 commit into `main` from `roles-feature`

Conversation 0 Commits 1 Checks 1 Files changed 34 +1,328 -0

obaigbenaa commented now Owner ...

No description provided.

commit new role files into Github 9bf4779

Add more commits by pushing to the `roles-feature` branch on `obaigbenaa/ansible-config-mgt`

Require approval from specific reviewers before merging
Branch protection rules ensure specific people approve pull requests before they're merged. Add rule ×

Continuous integration has not been set up
GitHub Actions and several other apps can be used to automatically catch bugs and enforce style.

This branch has no conflicts with the base branch
Merging can be performed automatically.

Merge pull request You can also [open this in GitHub Desktop](#) or view [command line instructions](#).

Write Preview H B I ≡ <> 🔗 📄 📁 🔍 👤 🔄 🔒

Leave a comment

Attach files by dragging & dropping, selecting or pasting them.

Close pull request Comment

Reviewers
No reviews
Still in progress? Convert to draft

Assignees
No one—assign yourself

Labels
None yet

Projects
None yet

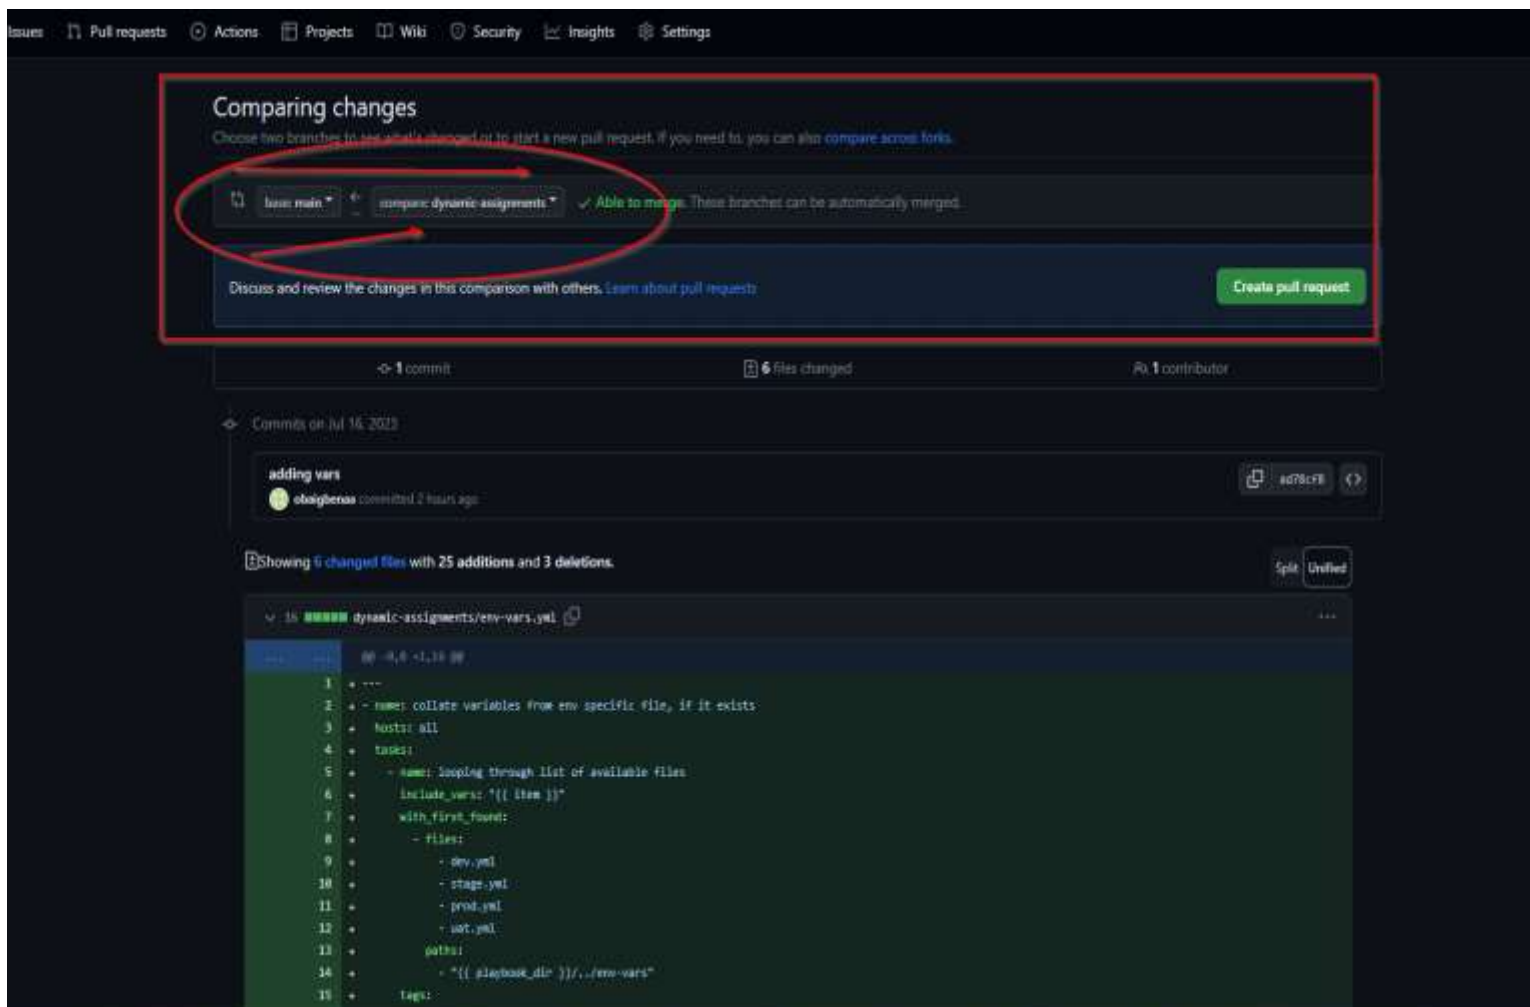
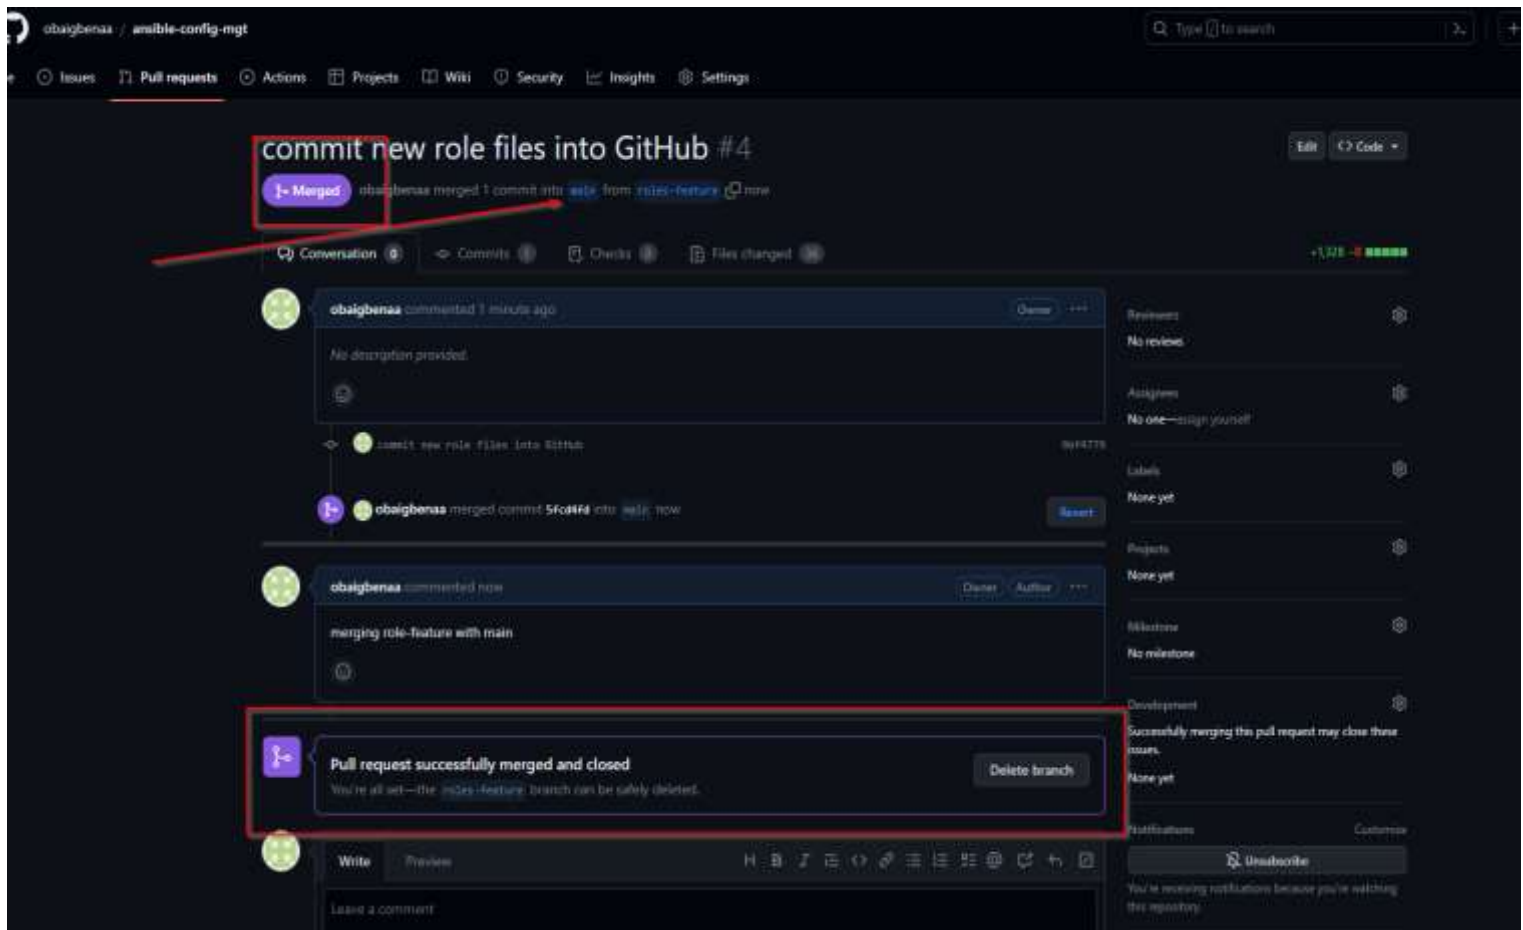
Milestone
No milestone

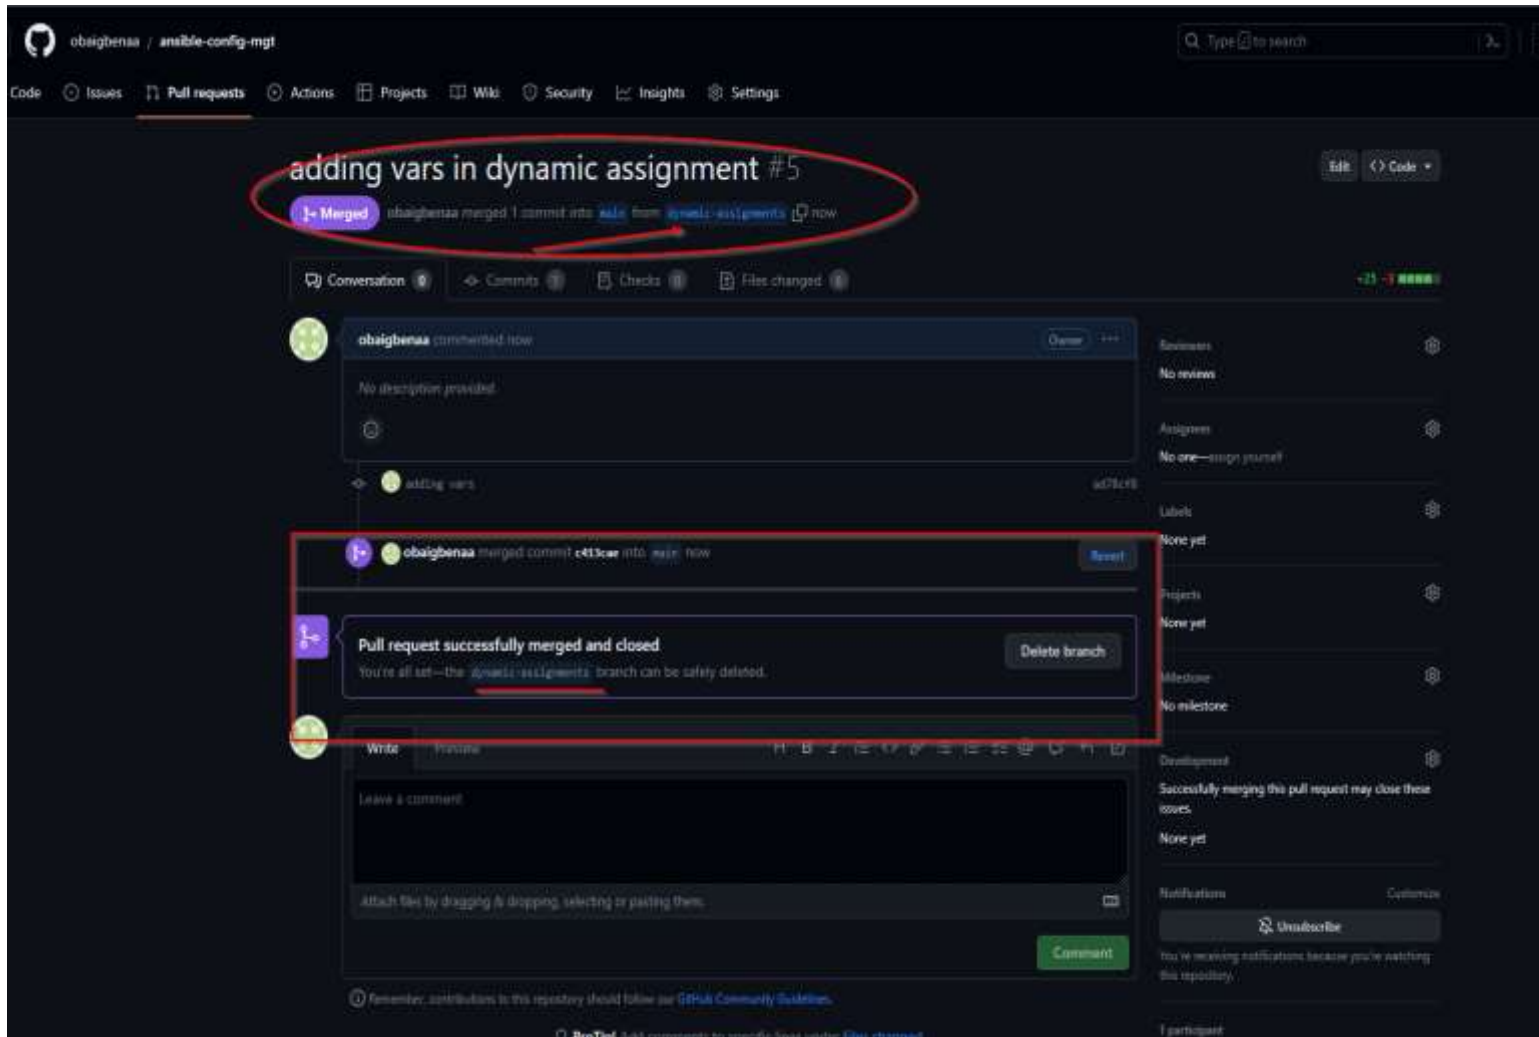
Development
Successfully merging this pull request may close these issues.
None yet

Notifications Customize
Unsubscribe
You're receiving notifications because you're watching this repository.

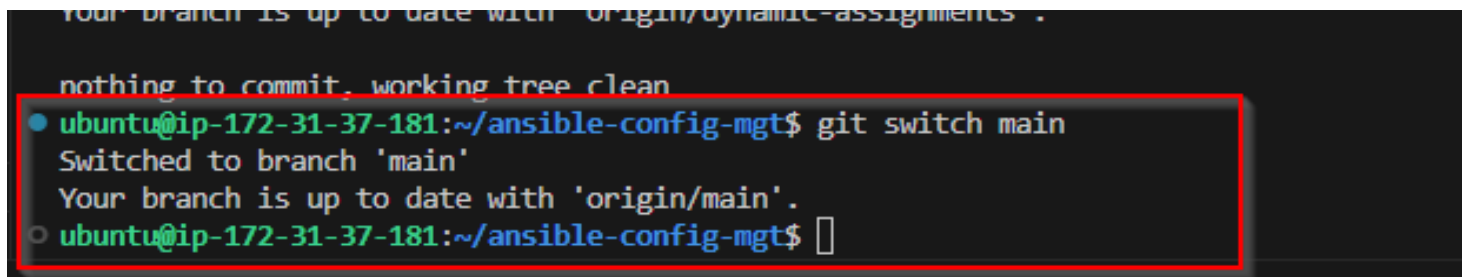
1 participant

Lock conversation





```
$ git switch main
```



LOAD BALANCER ROLES

We want to be able to choose which Load Balancer to use, **Nginx** or **Apache**, so we need to have two roles respectively:

1. Nginx
2. Apache

With your experience on Ansible so far you can:

- Decide if you want to develop your own roles, or find available ones from the community.

#I find it faster to develop roles from the community.

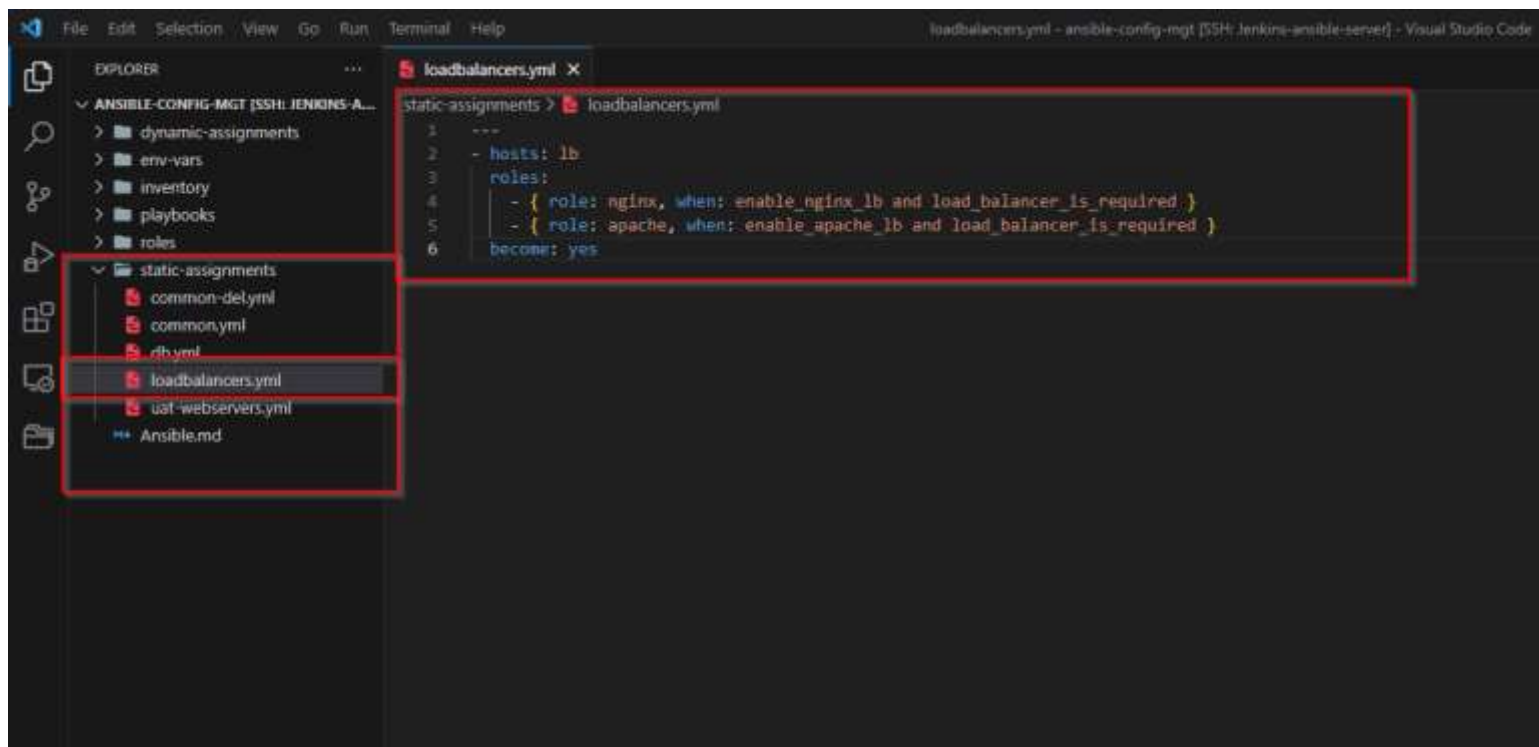
```
$ ansible-galaxy install geerlingguy.nginx
```

```
$ ansible-galaxy install geerlingguy.apache
```

```
$ mv geerlingguy.nginx/ nginx
```

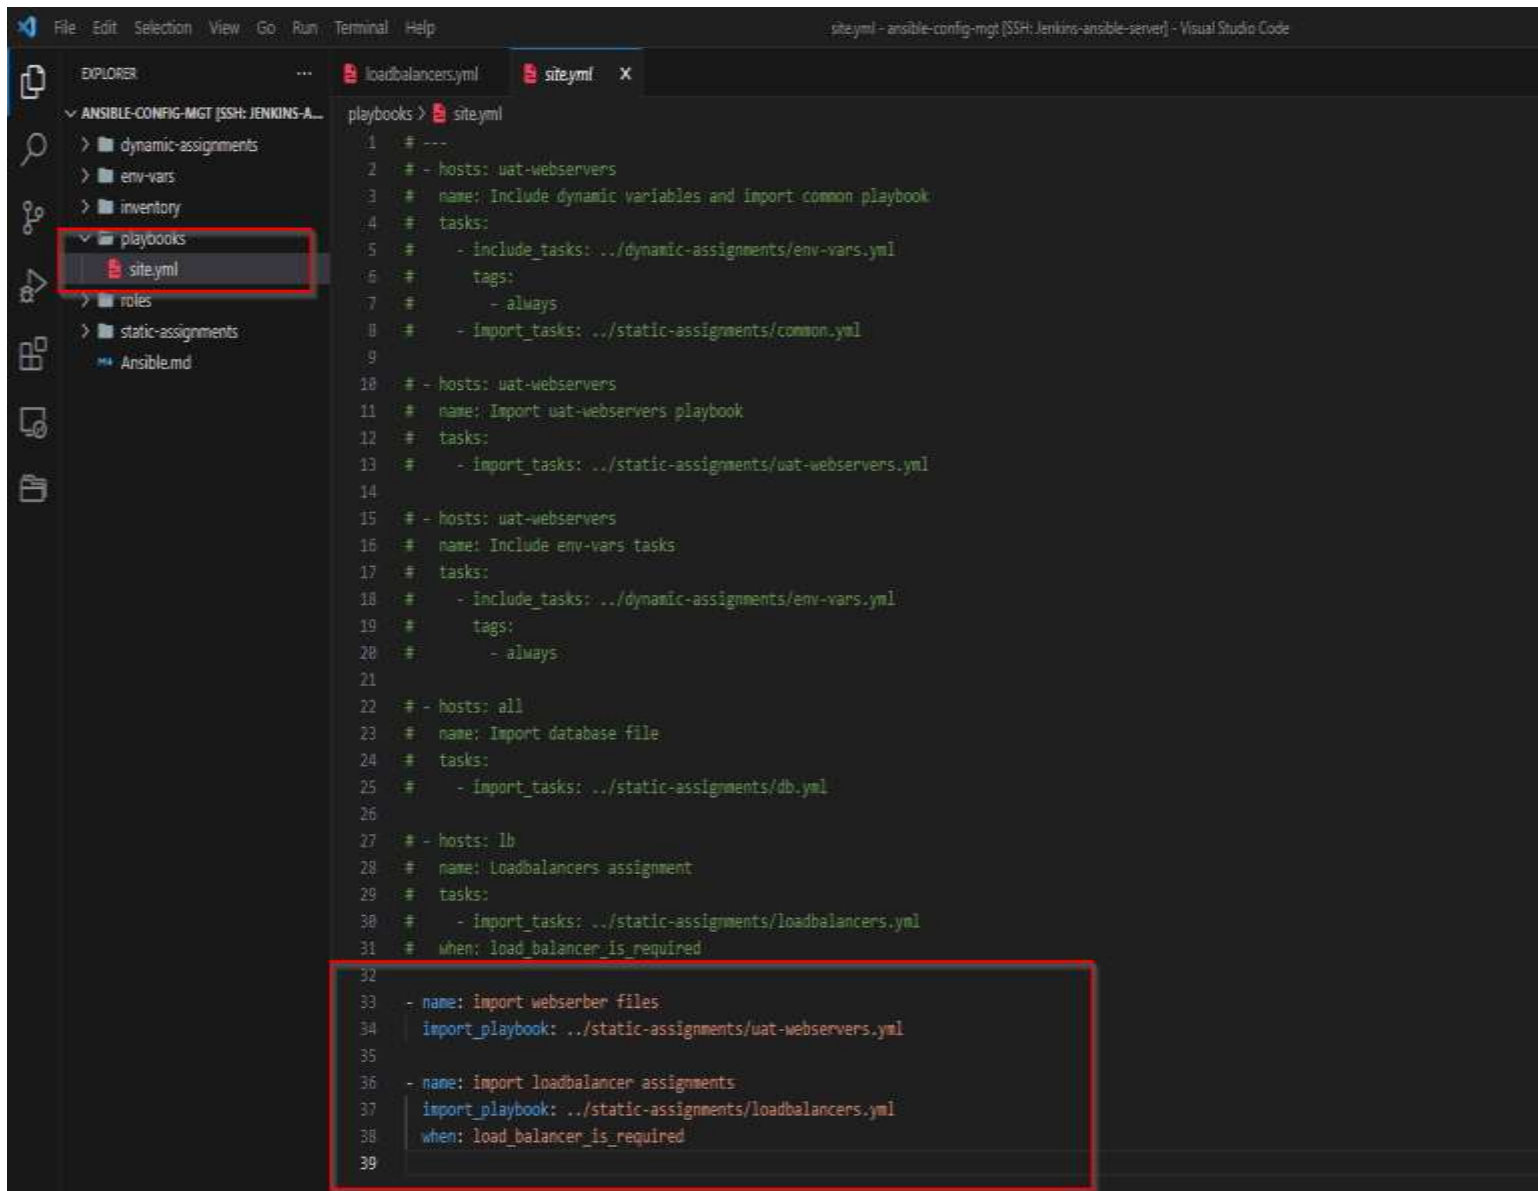
```
$ mv geerlingguy.apache/ apache
```

- Update both `static-assignment` and `site.yml` files to refer the roles



Important Hints:

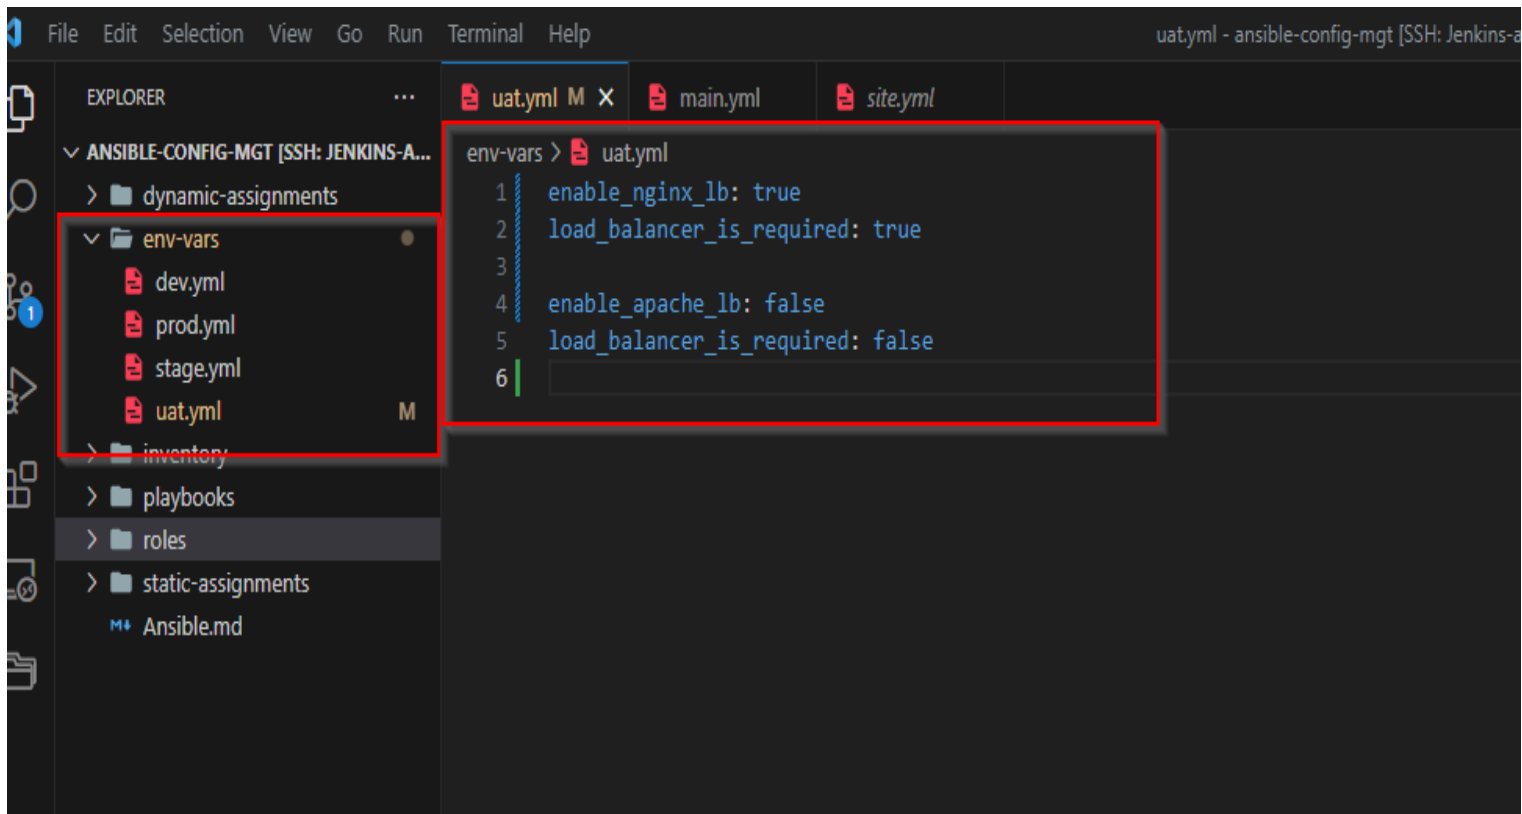
- Since we cannot use both **Nginx** and **Apache** load balancer, we need to add a condition to enable either one – this is where we can make use of variables.
- Declare a variable in `defaults/main.yml` file inside the Nginx and Apache roles. Name each variables `enable_nginx_lb` and `enable_apache_lb` respectively.
- Set both values to false like this `enable_nginx_lb: false` and `enable_apache_lb: false`.
- Declare another variable in both roles `load_balancer_is_required` and set its value to `false` as well.



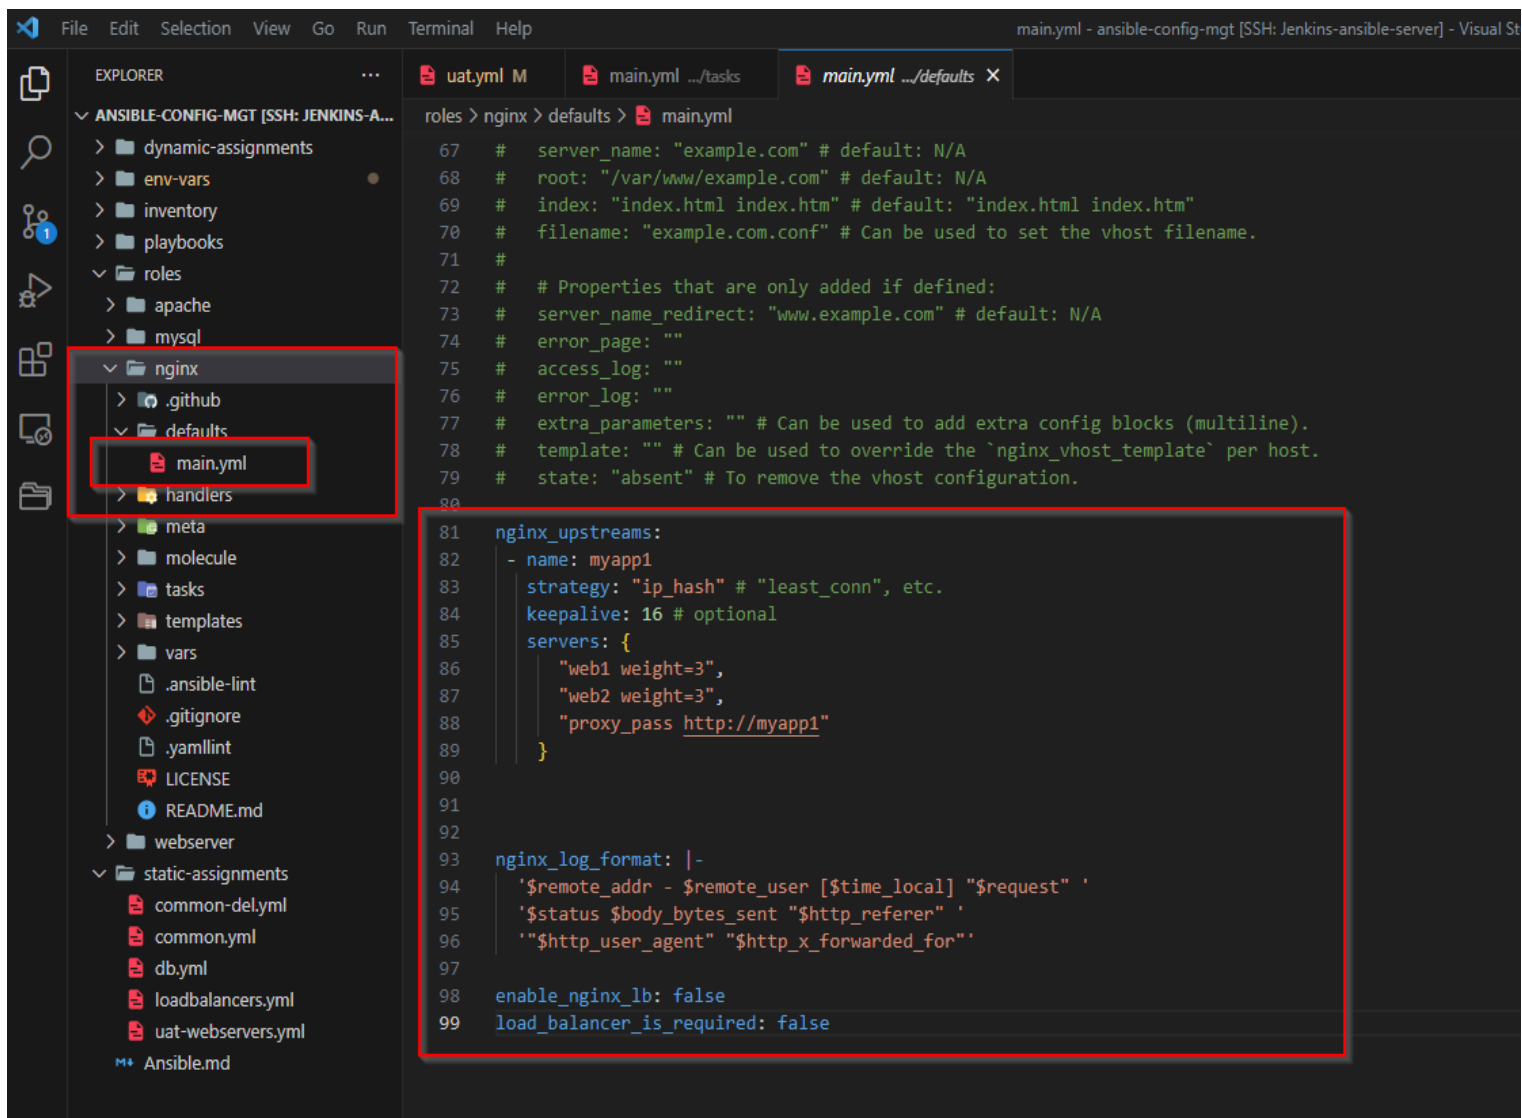
```
1 # ---
2 # - hosts: uat-webservers
3 #   name: Include dynamic variables and import common playbook
4 #   tasks:
5 #     - include_tasks: ../dynamic-assignments/env-vars.yml
6 #     tags:
7 #       - always
8 #     - import_tasks: ../static-assignments/common.yml
9
10 # - hosts: uat-webservers
11 #   name: Import uat-webservers playbook
12 #   tasks:
13 #     - import_tasks: ../static-assignments/uat-webservers.yml
14
15 # - hosts: uat-webservers
16 #   name: Include env-vars tasks
17 #   tasks:
18 #     - include_tasks: ../dynamic-assignments/env-vars.yml
19 #     tags:
20 #       - always
21
22 # - hosts: all
23 #   name: Import database file
24 #   tasks:
25 #     - import_tasks: ../static-assignments/db.yml
26
27 # - hosts: lb
28 #   name: Loadbalancers assignment
29 #   tasks:
30 #     - import_tasks: ../static-assignments/loadbalancers.yml
31 #   when: load_balancer_is_required
32
33 # - name: import webserver files
34 #   import_playbook: ../static-assignments/uat-webservers.yml
35
36 # - name: import loadbalancer assignments
37 #   import_playbook: ../static-assignments/loadbalancers.yml
38 #   when: load_balancer_is_required
39
```

Note the content of the file that has been commented out.

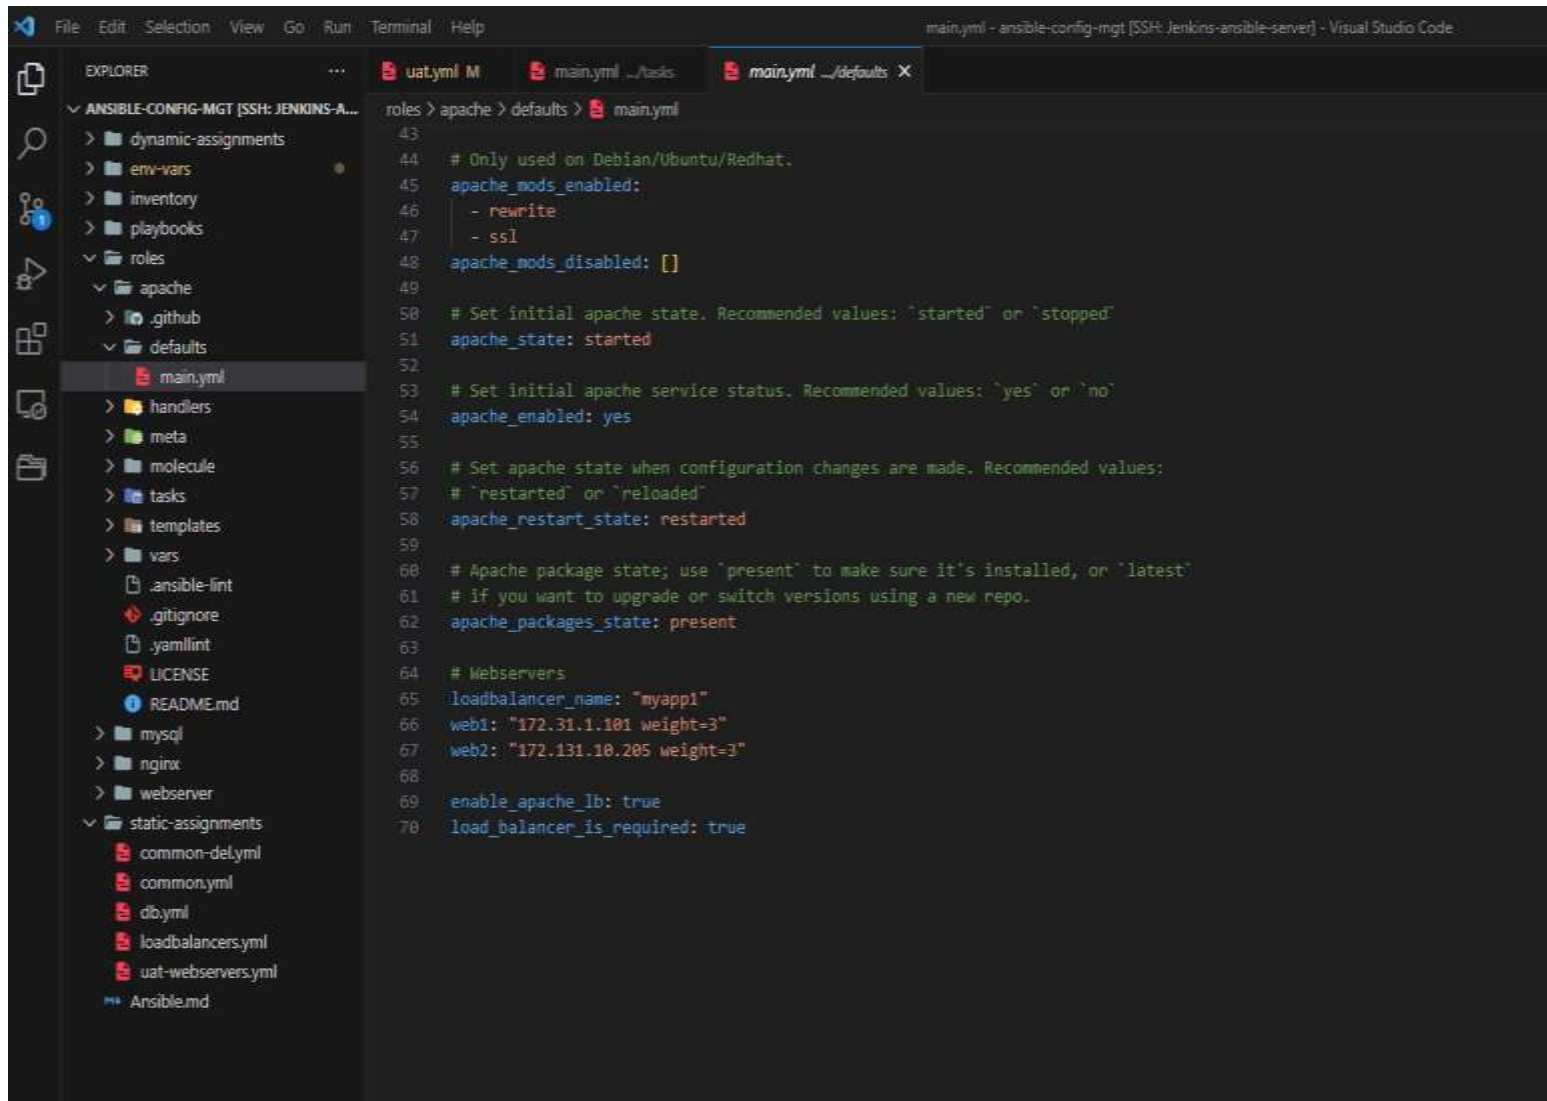
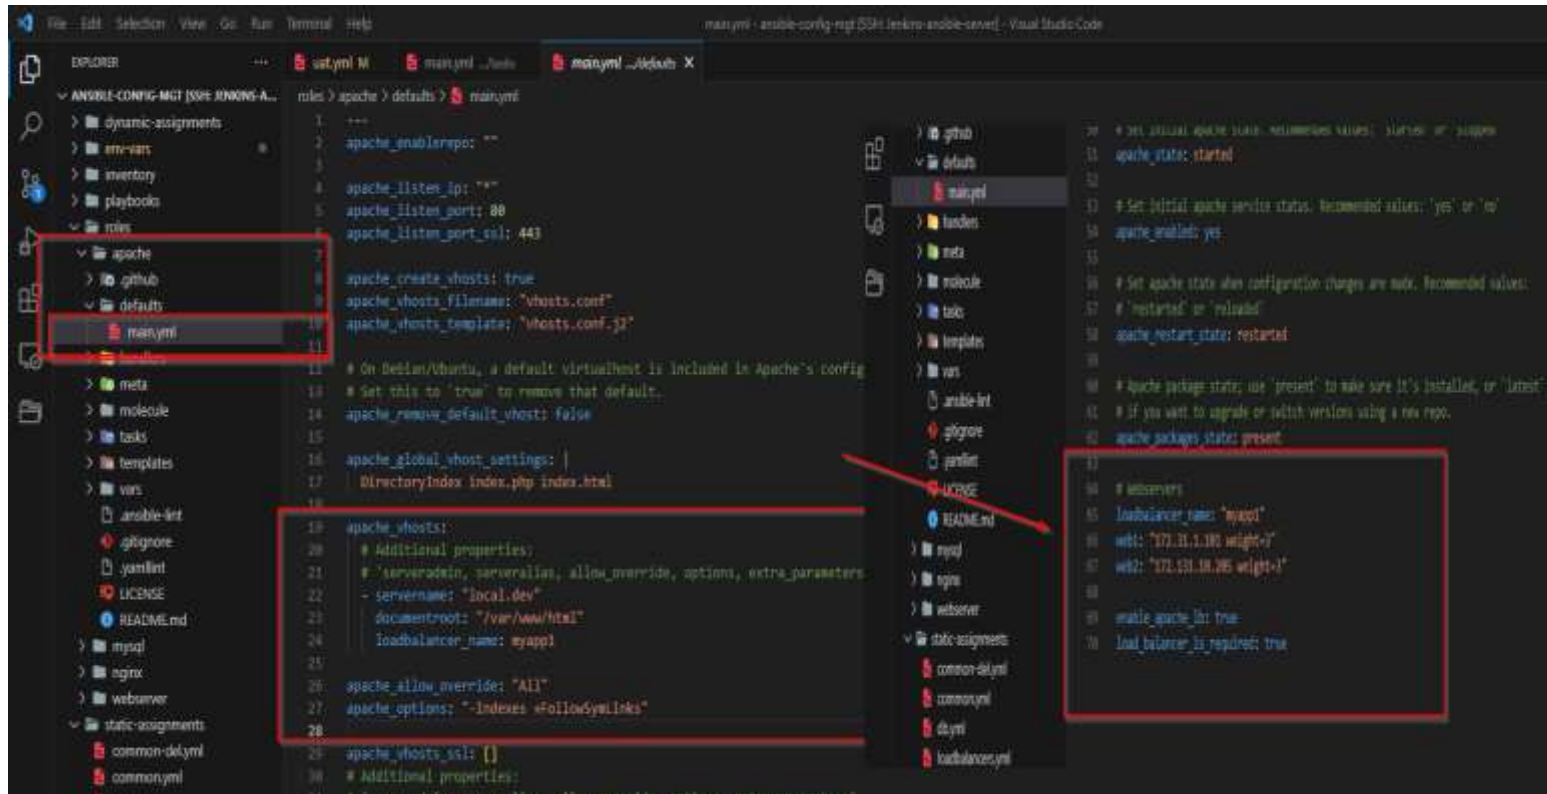
We can make use of **env-vars\uat.yml** file to define which loadbalancer to use in UAT environment by setting respective environmental variable to **true**.



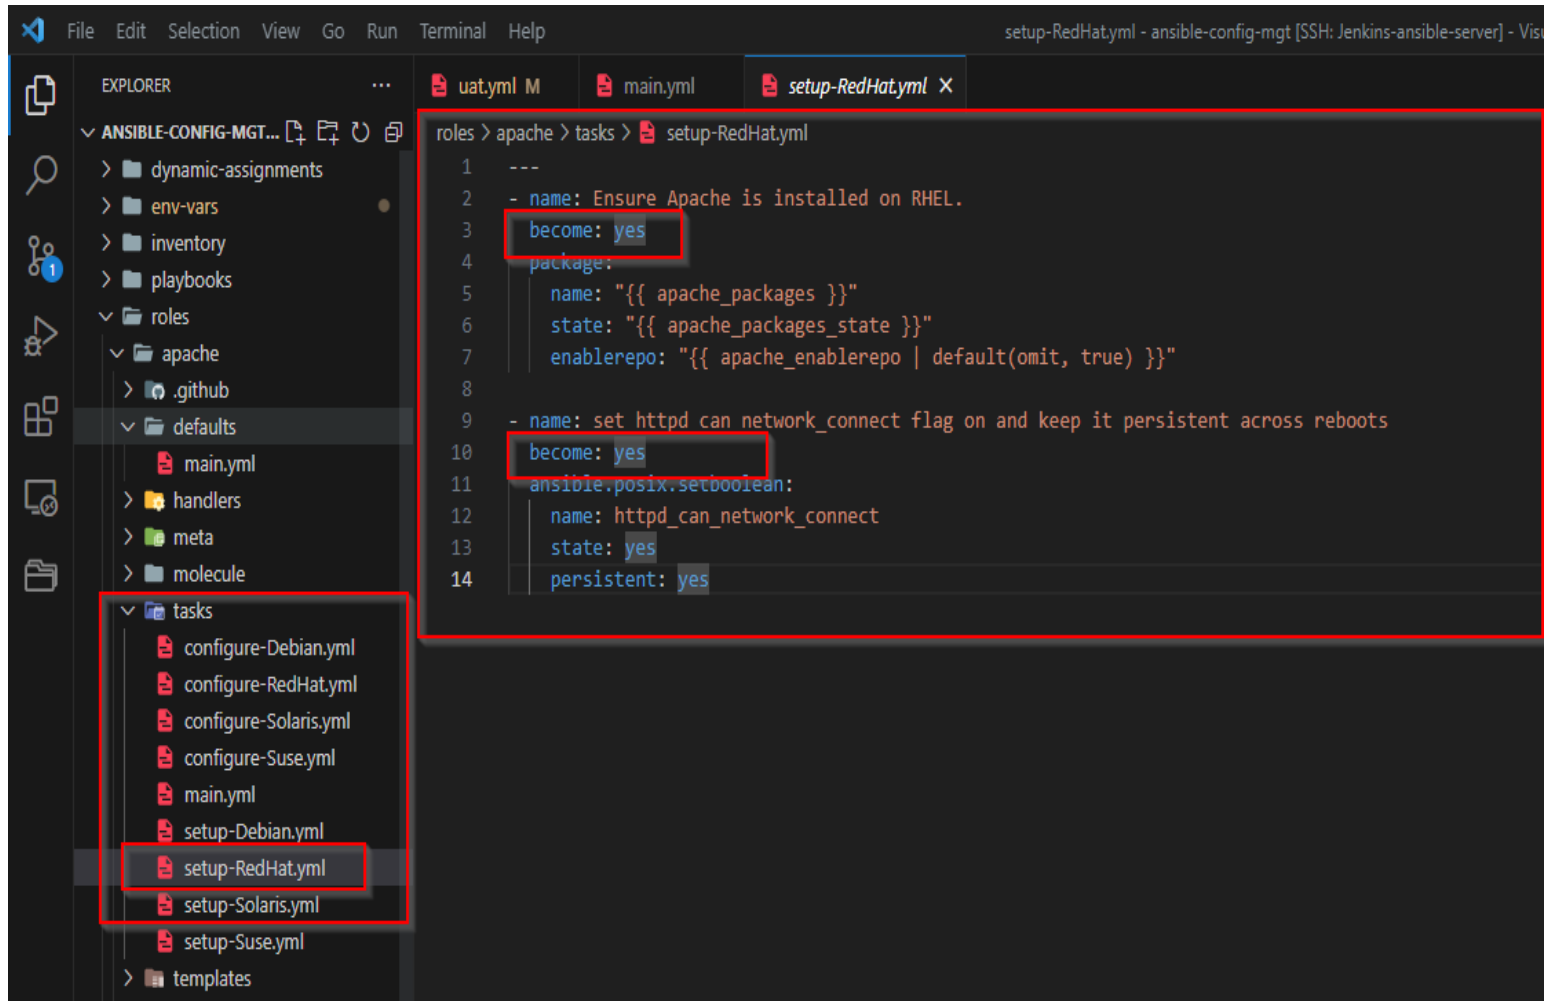
Configure nginx loadbalancer *defaults/main.yml*



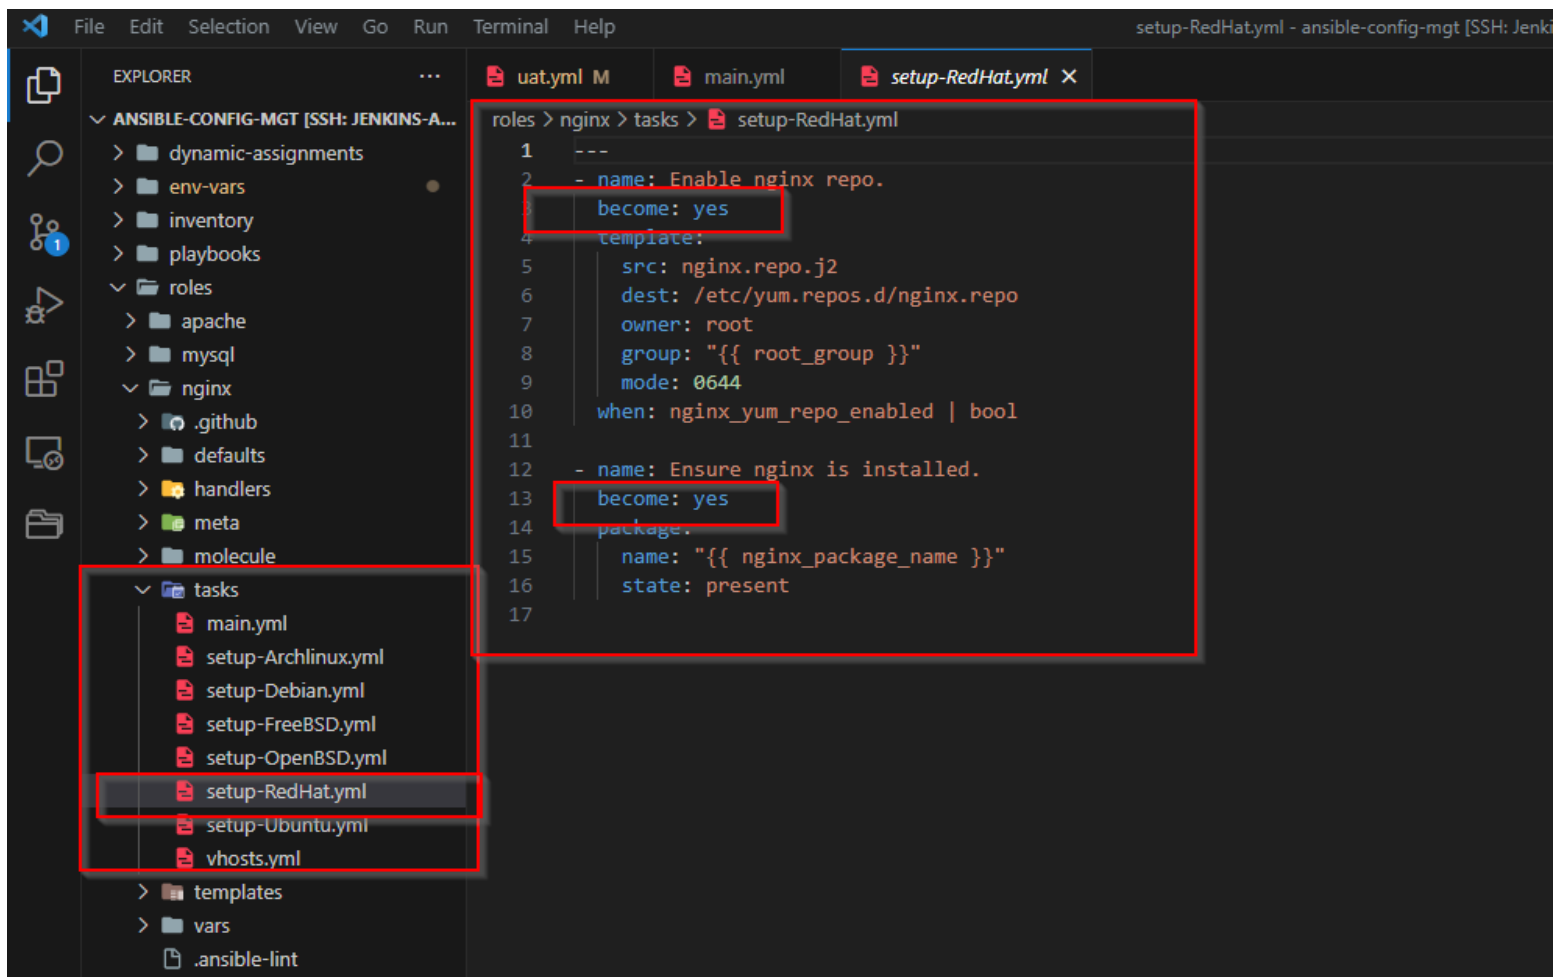
Configure apache loadbalancer *defaults/main.yml*



Enable super user on apache on *roles/apache/tasks/setup-RedHat.yml*



Enable super user on nginx on *roles/nginx/tasks/setup-RedHat.yml*



Run Ansible against the `uat` and load balancer environment.

```
$ ansible-playbook -i inventory/uat.yml playbooks/site.yml
```

```
PLAY RECAP *****
172.17.0.1:181      : ok=7  changed=3  unreachable=0  failed=0  skipped=0  rescued=0  ignored=0
172.17.0.20:181    : ok=7  changed=3  unreachable=0  failed=0  skipped=0  rescued=0  ignored=0
172.17.0.45:181    : ok=5   changed=0  unreachable=0  failed=0  skipped=1  rescued=0  ignored=0
```

Congratulations!!!

We have learned and practiced how to use Ansible configuration management tool to prepare UAT environment for Tooling web solution.